

DAYTONA HACKSPRINT SEOUL · SEPTEMBER 19, 2026

StrategyProof

Don't trust a trading strategy. Test it.

Describe a trading idea in plain English. StrategyProof turns it into precise rules, executes them on real market history inside an isolated Daytona sandbox, and has an independent open model on Nosana challenge the evidence.

Kimi · strategy parser

Daytona · isolated execution

Nosana · independent reviewer

THE PROBLEM

Finding trading strategies is easy. Knowing if they ever worked is not.

1,000s

of strategies are shared every day on YouTube, Telegram, Discord and X, usually as screenshots and claims.

3 walls

Validating one takes code, clean historical data and backtesting infrastructure. Most traders have none of them.

0 proof


So people trade on trust. Hype replaces evidence, and the cost shows up in real accounts.

Can we replace "trust this strategy" with "show me the evidence", in under 30 seconds?

THE SOLUTION

One sentence in. Reproducible evidence out.

- **Kimi** turns the idea into a strict, machine-readable strategy
- **Daytona** gives that strategy its own temporary computer
- It runs against **48,856 real EURUSD 15m candles**
- The sandbox is **destroyed** when the test finishes
- **Nosana** runs a second AI that attacks the results



Natural language → rules → isolated execution → critique



Key design choice


The LLM **never writes executable code**. It only fills in a small JSON DSL. The sandbox runs our own deterministic, stdlib-only engine: signals on candle close, fills at the next open, stop loss wins same-candle ties. User text never reaches a shell.

A real sandbox is created, used and destroyed on every run

AGENT ACTIVITY		26.4s
✓	Kimi converts natural language into trading rules RSI 14 Cross Above 30 Long · long · RSI(14) crosses above 30	5.3s
✓	Create isolated Daytona sandbox 0486f644-d77d-40b2-9403-7230bc17d277 · region eu	1.3s
✓	Upload backtest.py, strategy.json and EURUSD history 48,856 candles of real EURUSD 15m data	2.5s
✓	Execute historical simulation inside the sandbox 395 trades simulated · engine 162 ms · host 0486f644-d77d-40b2-9403-7230bc17d277	1.0s
✓	Destroy Daytona sandbox Sandbox 0486f644-d77... no longer exists	1.2s
✓	Nosana independently reviews the evidence open model qwen/qwen3.8-27b · 4 warnings raised	13.1s

THE EVIDENCE

"Buy EURUSD when RSI 14 crosses back above 30. 20 pip stop, 40 pip target."



A second AI attacks the evidence

An open model served on **Nosana** gets the formal strategy, the metrics and per-segment results. Its job is to find weaknesses, not to say "trade this".

- Negative expectancy, flagged
- Edge only in the first third of the data
- Spread eats 4% of risk per trade
- Concrete next tests to run

Robustness review

Independent · Nosana · qwen/qwen3.8-27b

The strategy exhibits a negative expectancy with a profit factor below 1.0 and a maximum drawdown significantly exceeding the net loss, indicating that the current RSI entry logic and fixed risk-reward ratio are not statistically robust for this instrument and timeframe.

HIGH ⚠ Negative Expectancy and Profit Factor

The strategy has a profit factor of 0.964 and an average R of -0.016, meaning it loses money on average. With a 34.43% win rate and a 2:1 reward-to-risk ratio, the break-even win rate is approximately 33.3%. The current performance is barely below break-even before accounting for slippage or additional costs, making it highly vulnerable to minor adverse market conditions.

HIGH ⚠ Excessive Drawdown Relative to Return

The maximum drawdown of 27.83% is nearly three times the net return of -9.77%. This indicates that the strategy's equity curve is highly volatile and that the capital required to sustain the strategy is disproportionately high compared to its potential (negative) output. The risk of ruin is significant if the trend continues.

MEDIUM ⚠ Performance Instability Across Segments

The strategy shows inconsistent performance across time segments: positive in the early period (+7.88%), but negative in both the middle (-10.51%) and recent (-6.55%) periods. This lack of consistency suggests the strategy may be sensitive to specific market regimes (e.g., trending vs. ranging) and lacks robustness across different volatility environments.

MEDIUM ⚠ Sensitivity to Transaction Costs

The strategy incurs a spread cost of 0.8 pips per trade. Given the tight stop loss (20 pips) and take profit (40 pips), the spread represents a significant portion of the risk (4%) and reward (2%). Any increase in spread or slippage would further degrade the already marginal profit factor, potentially pushing it well below 0.9.

WHAT HOLDS UP

- ✓ The sample size of 395 trades is sufficient to provide a statistically meaningful baseline for evaluating the strategy's core mechanics.
- ✓ The risk management parameters (1% risk per trade) are conservative, limiting the impact of any single trade on the overall equity curve.

EVIDENCE TO COLLECT NEXT

- Test the strategy with a wider stop loss (e.g., 30-40 pips) to see if the win rate improves sufficiently to offset the lower reward-to-risk ratio, or if the drawdown becomes more manageable.
- Analyze the performance of the strategy during high-volatility periods (e.g., NFP, CPI releases) to determine if the RSI signal is being triggered by noise rather than genuine momentum shifts.
- Compare the strategy's performance against a simple buy-and-hold or a random entry baseline to determine if the RSI filter provides any statistical edge over random timing in this specific market environment.
- Evaluate the impact of varying the RSI period (e.g., 7, 21) and threshold (e.g., 20, 40) to check if the current parameters are overfitted to the specific 2024-2025

BUILT ON THE SPONSORS, AT CODE LEVEL

Integrations

SPONSOR	CODE	WHAT IT DOES
Daytona	<code>lib/daytona.ts</code> <code>@daytona/sdk</code>	One fresh ephemeral Python sandbox per request: <code>daytona.create</code> → <code>fs.uploadFiles</code> (engine + strategy + 3 MB dataset) → <code>process.executeCommand</code> (fixed command, 60 s timeout) → <code>sandbox.delete(60, true)</code> in <i>finally</i> . Sandbox ID is shown live in the UI.
Nosana	<code>lib/nosana.ts</code> <code>inference.nosana.com/v1</code>	OpenAI-compatible client, model auto-discovery via <code>GET /models</code> (cached), JSON-mode review validated with Zod, one repair attempt. A review failure never throws away a valid backtest.
Kimi	<code>lib/kimi.ts</code>	Kimi K2.5 turns natural language into the Strategy DSL, with a single repair request at most. Unsupported ideas (e.g. "liquidity sweep into FVG") are rejected explicitly.

Next.js 16 · TypeScript · Zod

Python stdlib engine

Streaming NDJSON

8 tests: DSL limits · determinism · SL-first ties

No secrets client-side

VISION

Today we validate one idea.

Next: autonomous StrategyProof agents that continuously test the strategies shared online, replacing screenshots, claims and hype with **reproducible evidence**.

Ingest

Extract strategies from YouTube, Telegram and TradingView posts and chart images

Stress-test

Walk-forward, Monte Carlo, parameter-sensitivity maps and cross-market runs, each in parallel sandboxes

Publish

Public StrategyProof pages, creator verification and FXToolkit tick-level data plus MT4/MT5 validation